

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sb
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from sklearn import metrics
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: df = pd.read_csv('WineQTdataset.csv')
```

```
In [3]: df
```

Out[3]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56
...
1138	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75
1139	6.8	0.620	0.08	1.9	0.068	28.0	38.0	0.99651	3.42	0.82
1140	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58
1141	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76
1142	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71

1143 rows × 11 columns



```
In [4]: df.head(5)
```

Out[4]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alc
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	



```
In [5]: df.tail(5)
```

```
Out[5]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates
1138	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75
1139	6.8	0.620	0.08	1.9	0.068	28.0	38.0	0.99651	3.42	0.82
1140	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58
1141	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76
1142	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 1143 entries, 0 to 1142  
Data columns (total 13 columns):  
#   Column                Non-Null Count  Dtype  
---  -  
0   fixed acidity          1143 non-null   float64  
1   volatile acidity       1143 non-null   float64  
2   citric acid            1143 non-null   float64  
3   residual sugar         1143 non-null   float64  
4   chlorides              1143 non-null   float64  
5   free sulfur dioxide    1143 non-null   float64  
6   total sulfur dioxide   1143 non-null   float64  
7   density                1143 non-null   float64  
8   pH                    1143 non-null   float64  
9   sulphates              1143 non-null   float64  
10  alcohol                1143 non-null   float64  
11  quality                1143 non-null   int64  
12  Id                     1143 non-null   int64  
dtypes: float64(11), int64(2)  
memory usage: 116.2 KB
```

```
In [7]: df.isnull()
```

```
Out[7]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates
0	False	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False	False
...
1138	False	False	False	False	False	False	False	False	False	False
1139	False	False	False	False	False	False	False	False	False	False
1140	False	False	False	False	False	False	False	False	False	False
1141	False	False	False	False	False	False	False	False	False	False
1142	False	False	False	False	False	False	False	False	False	False

1143 rows × 13 columns



```
In [8]: df.sum()
```

```
Out[8]: fixed acidity      9499.600000
volatile acidity      607.320000
citric acid          306.740000
residual sugar      2894.250000
chlorides           99.364000
free sulfur dioxide  17848.500000
total sulfur dioxide  52480.500000
density            1139.262860
pH                 3784.490000
sulphates           751.760000
alcohol            11935.333333
quality             6466.000000
Id                920080.000000
dtype: float64
```

```
In [9]: df.isnull().sum()
```

```
Out[9]: fixed acidity      0
volatile acidity    0
citric acid         0
residual sugar      0
chlorides           0
free sulfur dioxide 0
total sulfur dioxide 0
density             0
pH                  0
sulphates           0
alcohol             0
quality             0
Id                  0
dtype: int64
```

```
In [10]: df.describe()
```

```
Out[10]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total d
count	1143.000000	1143.000000	1143.000000	1143.000000	1143.000000	1143.000000	1143.0
mean	8.311111	0.531339	0.268364	2.532152	0.086933	15.615486	45.9
std	1.747595	0.179633	0.196686	1.355917	0.047267	10.250486	32.7
min	4.600000	0.120000	0.000000	0.900000	0.012000	1.000000	6.0
25%	7.100000	0.392500	0.090000	1.900000	0.070000	7.000000	21.0
50%	7.900000	0.520000	0.250000	2.200000	0.079000	13.000000	37.0
75%	9.100000	0.640000	0.420000	2.600000	0.090000	21.000000	61.0
max	15.900000	1.580000	1.000000	15.500000	0.611000	68.000000	289.0

```
In [11]: df = pd.get_dummies(df,drop_first=True)
df
```

Out[11]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56
...
1138	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75
1139	6.8	0.620	0.08	1.9	0.068	28.0	38.0	0.99651	3.42	0.82
1140	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58
1141	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76
1142	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71

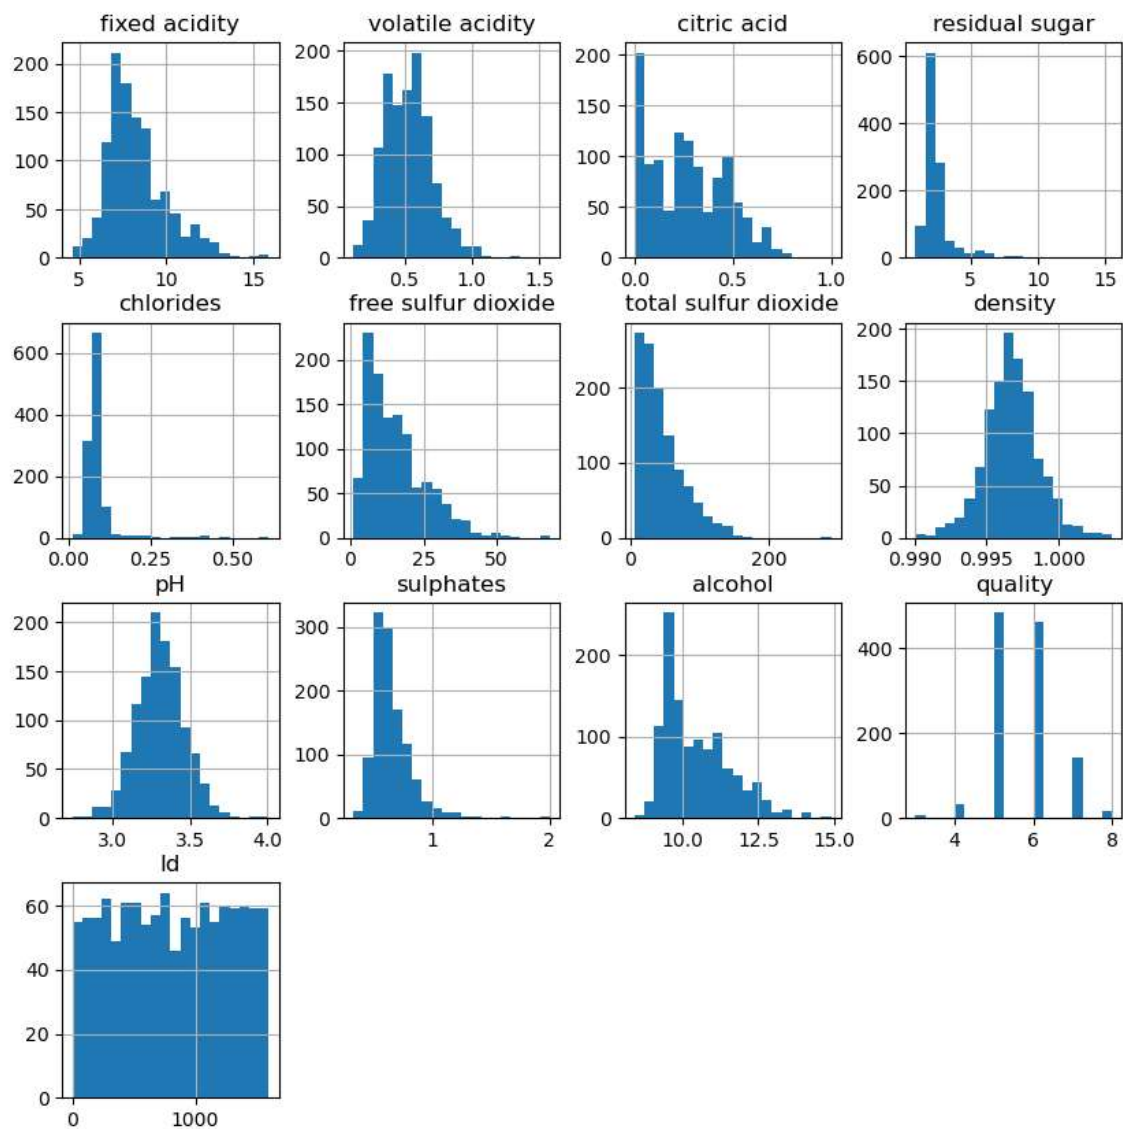
1143 rows × 13 columns




```
In [13]: for col in df.columns:
          if df[col].isnull().sum() > 0:
              df[col] = df[col].fillna(df[col].mean())
df.isnull().sum().sum()
```

Out[13]: 0


```
In [14]: df.hist(bins=20, figsize=(10, 10))
plt.show()
```



```
In [15]: plt.bar(df['quality'], df['alcohol'])  
plt.xlabel('quality')  
plt.ylabel('alcohol')  
plt.show()
```



```
In [18]: plt.figure(figsize=(8, 8))
sb.heatmap(df.corr() > 0.7, annot=True, cbar=False)
plt.show()
```



```
In [19]: df = df.drop('total sulfur dioxide', axis=1)
```

```
In [20]: df['best quality'] = [1 if x > 5 else 0 for x in df.quality]
df.replace({'white': 1, 'red': 0}, inplace=True)
```

```
In [21]: features = df.drop(['quality', 'best quality'], axis=1)
target = df['best quality']
xtrain, xtest, ytrain, ytest = train_test_split(
    features, target, test_size=0.2, random_state=40)
xtrain.shape, xtest.shape
```

```
Out[21]: ((914, 11), (229, 11))
```


In [22]: xtrain

Out[22]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	density	pH	sulphates	alcohol
675	9.0	0.36	0.52	2.10	0.111	5.0	0.99568	3.31	0.62	11.3
653	8.6	0.47	0.27	2.30	0.055	14.0	0.99516	3.18	0.80	11.2
845	7.7	0.57	0.21	1.50	0.069	4.0	0.99458	3.16	0.54	9.8
1027	6.9	0.58	0.20	1.75	0.058	8.0	0.99322	3.38	0.49	11.7
1023	10.0	0.38	0.38	1.60	0.169	27.0	0.99914	3.15	0.65	8.5
...
626	6.8	0.57	0.00	2.50	0.072	32.0	0.99491	3.43	0.56	11.2
1016	7.6	0.41	0.33	2.50	0.078	6.0	0.99570	3.30	0.58	11.2
165	8.2	1.00	0.09	2.30	0.065	7.0	0.99685	3.32	0.55	9.0
7	7.3	0.65	0.00	1.20	0.065	15.0	0.99460	3.39	0.47	10.0
219	8.4	0.65	0.60	2.10	0.112	12.0	0.99730	3.20	0.52	9.2

914 rows × 11 columns



In [23]: xtest

Out[23]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	density	pH	sulphates	alcohol
471	7.2	0.57	0.06	1.6	0.076	9.0	0.99720	3.36	0.70	9.6
192	11.5	0.18	0.51	4.0	0.104	4.0	0.99960	3.28	0.97	10.1
1035	6.5	0.90	0.00	1.6	0.052	9.0	0.99467	3.50	0.63	10.9
476	8.2	0.73	0.21	1.7	0.074	5.0	0.99680	3.20	0.52	9.5
512	8.4	0.56	0.04	2.0	0.082	10.0	0.99760	3.22	0.44	9.6
...
281	7.7	0.69	0.05	2.7	0.075	15.0	0.99740	3.26	0.61	9.1
176	7.1	0.60	0.00	1.8	0.074	16.0	0.99720	3.47	0.70	9.9
537	8.3	0.65	0.10	2.9	0.089	17.0	0.99803	3.29	0.55	9.5
1043	7.3	0.48	0.32	2.1	0.062	31.0	0.99728	3.30	0.65	10.0
801	8.5	0.28	0.35	1.7	0.061	6.0	0.99524	3.30	0.74	11.8

229 rows × 11 columns



```
In [24]: ytrain
```

```
Out[24]: 675      1
          653      0
          845      1
          1027     0
          1023     0
          ..
          626      1
          1016     0
          165      1
           7       1
          219      0
          Name: best quality, Length: 914, dtype: int64
```

```
In [25]: ytest
```

```
Out[25]: 471      1
          192      1
          1035     1
          476      0
          512      0
          ..
          281      0
          176      1
          537      0
          1043     1
          801      1
          Name: best quality, Length: 229, dtype: int64
```

```
In [26]: xtrain.shape
```

```
Out[26]: (914, 11)
```

```
In [27]: xtest.shape
```

```
Out[27]: (229, 11)
```

```
In [30]: ytrain.shape
```

```
Out[30]: (914,)
```

```
In [29]: ytest.shape
```

```
Out[29]: (229,)
```

```
In [31]: norm = MinMaxScaler()
          xtrain = norm.fit_transform(xtrain)
          xtest = norm.transform(xtest)
```

```
In [46]: models = [LogisticRegression(),SVC(kernel='rbf')]
for i in range(2):
    models[i].fit(xtrain, ytrain)
    print(f'models[i] : ')
    print('Training Accuracy : ', metrics.roc_auc_score(ytrain, models[i].predict(xtrain)))
    print('Validation Accuracy : ', metrics.roc_auc_score(ytest, models[i].predict(xtest)))
    print()
```

```
models[i] :
Training Accuracy : 0.7546950559364851
Validation Accuracy : 0.7255154639175256
```

```
models[i] :
Training Accuracy : 0.7648213641284736
Validation Accuracy : 0.7358247422680412
```

```
In [36]: from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt

y_pred = models[1].predict(xtest)
cm = confusion_matrix(ytest, y_pred)
print(cm)
```

```
[[70 27]
 [33 99]]
```

```
In [38]: print(metrics.classification_report(ytest,
models[1].predict(xtest)))
```

	precision	recall	f1-score	support
0	0.68	0.72	0.70	97
1	0.79	0.75	0.77	132
accuracy			0.74	229
macro avg	0.73	0.74	0.73	229
weighted avg	0.74	0.74	0.74	229

```
In [43]: for a in range(len(df.corr().columns)):
for b in range(a):
    if abs(df.corr().iloc[a,b]) >0.7:
        name = df.corr().columns[a]
        print(name)
```

```
best quality
```

```
In [ ]:
```

